

CS 3753 & 5163 Data Science Fall 2020

Homework 4 (80 points)

Submission:

1. Submit a single python script (`abc123_hw#.ipynb` or `abc123_hw#.py`) through Blackboard learn. All the results are outputted from your Python code.
2. You should have **the instruction of running your code at the beginning of your code**. It should run successfully either in the basic Python3 environment or in Jupyter Notebook.
3. **If your code cannot run**, we assume your code can run, then we will check whether your code is correct logically. If so, half points will be deducted. Otherwise, more points will be deducted if your code is wrong or there is no code.
4. Do not compress your source code and data files. Make sure all your files are in the same folder when you run the code. So, after the graders download your homework, they do not need to set the path for the data file. They can run your code successfully.
5. If there is any plagiarism, you will lose all points on the questions at first time. In next, you will lose all points in the whole homework.
6. You can submit your homework **3 times** before the deadline. The late submission will lose **15%** of the total points in the assignment. The compressed files will receive a warning at the first time and will lose **10% points later**.

Questions


1. Linear Regression (20 points)

If the plot of n pairs of data (x, y) for an experiment appear to indicate a "linear relationship" between y and x , then the method of least squares may be used to write a linear relationship between x and y .

The least squares regression line is the line that minimizes the sum of the squares of the errors (SSE) from each data point to the line (see figure below). The least square regression line for the set of n data points is given by $y = \beta x + \alpha$, where α and β are given by

$$\beta = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}$$

$$\alpha = \bar{y} - \beta * \bar{x}$$



The sales of a company (in million dollars) for each year are shown in the table. Output your results by the print() function in Python.


x(year)	2005	2006	2007	2008	2009
y(sales)	12	19	29	37	45

- Write a Python code to find the least square regression line $y = \beta x + \alpha$. You either can write a Python code to calculate α and β based on the equations above or use the linear regression function in the model sklearn (described in the lecture) to get the results directly. (15 pts)
- Use the regression line as a model to estimate the sales of the company in 2012 (in million dollars). (5 pts)


2. Statistics (60 points)

Write a program to simulate an experiment of tossing a fair coin 16 times and counting the number of heads. Repeat this experiment 10^5 times to obtain the number of heads for every 16 tosses; save the number of heads in a vector of size 10^5 (nHeads). You should be able to do this in just a few lines. (Use np.random.uniform to generate a 2d array of $10^5 * 16$ random numbers between 0 and 1; a value that is greater than 0.5 is considered a “head”.) Complete the following questions.


- Plot the histogram of nHeads using plt.hist, with parameter bins = range(18). Label your plots clearly (10pts).



- b. Plot the PMF using `plt.hist` with parameter `bins` and `density`. Label your plots clearly (10pts).



- c. Calculate the probability of having NO MORE THAN k heads out of 16 tosses, where $k = 0, 1, 2, \dots, 15, 16$. Plot this as a CDF. You can calculate the probabilities/counts with your own code or using values returned from 2a/2b.) Label your plots clearly (10pts).



- d. Use the binomial distribution CDF (use `scipy.stats.binom.cdf`) to compute the probability of having NO MORE THAN k heads out of 16 tosses, where $k = 0, 1, 2, \dots, 15, 16$ and compare these probabilities with the probabilities you obtained in 2c. (Plot the probabilities you obtained from the simulation results in 2c against the probabilities from your theoretical calculation here, as **a scatter plot and line graph**. Plot in **loglog scale** to visualize small probabilities.) (30pts)

